

MSiA400_Lab2

Data Import

```
dat <- read_csv("gradAdmit.csv")

## Rows: 400 Columns: 4

## -- Column specification -----
## Delimiter: ","
## dbl (4): admit, gre, gpa, rank

##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

Problem 1a

```
set.seed(123)
n <- nrow(dat)
train_i <- sample.int(n = n, size = floor(n * 0.8), replace = FALSE)
train <- dat[train_i, ]
test <- dat[-train_i, ]
```

Problem 1b

Kernel Selection and Degree, gamma, and coef0 search

```
set.seed(123)
folds <- createFolds(1:nrow(train), k=5)
kernels <- c("linear", "polynomial", "radial", "sigmoid")
train_acc_by_fold <- numeric(5)
val_acc_by_fold <- numeric(5)
avg_train_acc <- numeric(4)
avg_val_acc <- numeric(4)

degrees <- c(2,3,4)
gammas <- c(0.01, 0.1, 1)
coef0s <- c(0.01, 0.1, 1, 10)
costs <- c(0.01, 0.01, 1, 10, 100)
acc_df <- tibble(
```

```

degree = numeric(),
gamma = numeric(),
coef0 = numeric(),
cost = numeric(),
kernel = character(),
avg_training_acc = numeric(),
avg_validation_acc = numeric()
)

for (k in seq_len(length(degrees))) {
  for (m in seq_len(length(gammas))) {
    for (n in seq_len(length(coef0s))) {
      for (p in seq_len(length(costs))) {
        for (i in seq_len(4)) {
          for (j in seq_len(5)) {
            curr_val_set <- train[folds[[j]], ]
            curr_train_set <- train[-folds[[j]], ]
            curr_svm <- suppressWarnings(svm(factor(admit) ~ ., data = curr_train_set,
              kernel = kernels[i],
              degree = degrees[k],
              gamma = gammas[m],
              coef0 = coef0s[n],
              cost = costs[p]))

            curr_train_pred <- predict(curr_svm, newdata = curr_train_set, type = 'response')
            curr_train_acc <- sum(curr_train_set$admit == curr_train_pred) / nrow(curr_train_set)

            curr_pred <- predict(curr_svm, newdata = curr_val_set, type = 'response')
            curr_val_acc <- sum(curr_val_set$admit == curr_pred) / nrow(curr_val_set)

            train_acc_by_fold[j] <- curr_train_acc
            val_acc_by_fold[j] <- curr_val_acc
          }
          acc_df <- acc_df %>% add_row(degree = degrees[k], gamma = gammas[m], coef0 = coef0s[n], cost = costs[p],
            train_acc = train_acc_by_fold, val_acc = val_acc_by_fold)
        }
      }
    }
  }
}

acc_df

```

```

## # A tibble: 720 x 7
##   degree gamma coef0 cost kernel avg_training_acc avg_validation_acc
##   <dbl> <dbl> <dbl> <dbl> <chr>          <dbl>          <dbl>
## 1     2    0.01  0.01  0.01 linear          0.675          0.675
## 2     2    0.01  0.01  0.01 polynomial    0.675          0.675
## 3     2    0.01  0.01  0.01 radial      0.675          0.675
## 4     2    0.01  0.01  0.01 sigmoid    0.675          0.675
## 5     2    0.01  0.01  0.01 linear      0.675          0.675
## 6     2    0.01  0.01  0.01 polynomial 0.675          0.675

```

```
## 7      2 0.01 0.01 0.01 radial          0.675          0.675
## 8      2 0.01 0.01 0.01 sigmoid        0.675          0.675
## 9      2 0.01 0.01 1 linear           0.685          0.669
## 10     2 0.01 0.01 1 polynomial        0.675          0.675
## # ... with 710 more rows
```

```
newdf <- acc_df[order(-acc_df$avg_validation_acc, -acc_df$avg_training_acc),]
newdf
```

```
## # A tibble: 720 x 7
##   degree gamma coef0 cost kernel avg_training_acc avg_validation_acc
##   <dbl> <dbl> <dbl> <dbl> <chr>          <dbl>          <dbl>
## 1      4 0.01 10     100 polynomial    0.725          0.694
## 2      4 0.01 10      10 polynomial    0.710          0.694
## 3      3 0.1  0.01 100 polynomial    0.718          0.691
## 4      3 1    0.01 1 polynomial    0.718          0.691
## 5      2 0.1  0.01 100 polynomial    0.710          0.691
## 6      2 0.1  1     10 polynomial    0.710          0.691
## 7      2 1    0.01 10 polynomial    0.710          0.691
## 8      2 1    0.1  1 polynomial    0.710          0.691
## 9      3 0.01 10     100 polynomial    0.710          0.691
## 10     4 0.01 10      1 polynomial    0.710          0.691
## # ... with 710 more rows
```

As seen from the dataframe above, polynomial kernel with degree 4, gamma 0.01, coef0 = 10, cost = 10 is most optimal. Both training and validation accuracy is high, and the difference between them is small, which indicates less overfitting.

problem 1c

```
best_m <- svm(factor(admit) ~ ., data = train,
               kernel = "polynomial",
               degree = 4,
               gamma = 0.01,
               coef0 = 10,
               cost = 10)
pred <- predict(best_m, newdata = test, type = 'response')
sum(test$admit == pred) / nrow(test)
```

```
## [1] 0.7375
```